

CSC2710 Analysis of Algorithms — Fall 2022

Unit 13 - Randomized Algorithms

13.1 Randomized Algorithms

- Las Vegas algorithms are randomized algorithms that always provide the correct solution, but get there through non-deterministic means
 - We will use this method to analyze random quicksort, or **RandQS**.
- Monte Carlo algorithms are such that they do not always provide the correct solution, but a solution with a probability of being correct.
 - With repetitions of the algorithm, the “correctness” of the solution approaches 100%.
 - We will use this method to solve the MINCUT problem.

13.1.1 Remember the Probability Rules

- For events ε_1 and ε_2 , they are said to be *independent* if and only if
$$\mathbf{Pr}[\varepsilon_1 \cap \varepsilon_2] = \mathbf{Pr}[\varepsilon_1] \times \mathbf{Pr}[\varepsilon_2]$$
- More generally, if ε_1 and ε_2 are not necessarily independent,
$$\mathbf{Pr}[\varepsilon_1 \cap \varepsilon_2] = \mathbf{Pr}[\varepsilon_1 | \varepsilon_2] \times \mathbf{Pr}[\varepsilon_2] = \mathbf{Pr}[\varepsilon_2 | \varepsilon_1] \times \mathbf{Pr}[\varepsilon_1]$$
- In general:
$$\mathbf{Pr}[\cap_{i=1}^n \varepsilon_i] = \mathbf{Pr}[\varepsilon_1] \times \mathbf{Pr}[\varepsilon_2 | \varepsilon_1] \times \mathbf{Pr}[\varepsilon_3 | \varepsilon_1 \cap \varepsilon_2] \dots \mathbf{Pr}[\varepsilon_k | \cap_{i=1}^{k-1} \varepsilon_i]$$

13.2 Analysis of Randomized QuickSort (Las Vegas)

- Recall that the QUICKSORT algorithm requires determining a pivot element at each step.
 - In the variation we discussed in class, we said that this would be the first element every step.
 - However, you can randomly select this pivot. This variation we will refer to as RANDQS, and it belongs to the family of *randomized algorithms*.

```
# Input: A subarray of A defined by indices 'left' and 'right'
# Output: The subarray is in sorted increasing order
def RandQS( A, left, right ):
    if( left < right ):                # Line 1
        # Partition using random pivot
        s = RandomPartition( A[ left : right ] )    # Line 2
        # Sort the lower and upper halves separately
        RandQS( A[ left : s-1 ] )                  # Line 3
        RandQS( A[ s+1 : right ] )                  # Line 4
    return A                                       # Line 5
```

- We did not do in-depth analysis of QuickSort earlier in the semester¹, but now we will introduce some tools that will allow us to examine its performance.
 - We want to analyze the *expected number of comparison* during the partitioning step of the algorithm.
 - Let A_i denote the element of *rank* i , or the i^{th} smallest element of A , for $1 \leq i \leq n$.
 - We will use a *random variable* X to denote comparisons.
 - Let X_{ij} be equal to 1 if elements A_i and A_j are compared during the execution of RandQS, and 0 otherwise.
 - Practice:** Can two elements ever be compared more than once?
 - Thus, a count of the *total number of comparisons* in RandQS is given by:

$$\sum_{i=1}^n \sum_{j>i} X_{ij}$$
 - We are interested in the *expected number of comparison* for RandQS:

$$E \left[\sum_{i=1}^n \sum_{j>i} X_{ij} \right] = \sum_{i=1}^n \sum_{j>i} E[X_{ij}]$$
 - Practice:** What is the *expected value* of an expression?
 - Let p_{ij} denote the probability that $X_{ij} = 1$. Since X_{ij} can only be 1 or 0,

$$E[X_{ij}] = p_{ij} \times 1 + (1 - p_{ij}) \times 0 = p_{ij}$$
 - Practice:** Do you see why?
 - This means we can calculate the number of comparisons once we know the probability that A_i and A_j are compared during an execution of RandQS. How can we accomplish this?

¹The hand-wave method!

13.2.1 Analyzing RandQS Using BST

- Treat a single execution of **RandQS** as a binary tree T .
 - Each node A_i of the tree represents a distinct element of A chosen as the pivot for one partitioning.
 - The node's left subtree are the elements less than A_i (partitioned to the left) and the right subtree are the elements greater than A_i , so the result of the partitioning.
 - * This is a binary search tree.
 - * A_i is compared with every element of each of its subtrees, but elements of each subtree are not compared with each other at the i^{th} iteration of the algorithm.
 - **Practice:** Do you see why?
 - * This implies that elements A_i and A_j are compared if and only if one is an ancestor of the other.
 - Let π be the *level-order traversal* of T , or the order of the chosen pivots. We can make the following observations:
 1. A_i and A_j are compared if and only if A_i or A_j occurs earlier in π than *any* element A_k such that $i < k < j$.
 - * Otherwise, elements A_i and A_j will be in opposite subtrees. However, if either element A_i or A_j is selected first, then they are compared.
 - * So we need to know the probability that A_i or A_j is selected first of all the elements between A_i and A_j
 2. Any of the elements A_i, A_{i+1}, \dots, A_j is *equally likely* to be the first chosen to be a pivot, and therefore the first to appear in the sequence π . Thus, the probability of the first being chosen being either A_i or A_j is $\frac{2}{j-i+1}$.
- Thus, $p_{ij} = \frac{2}{j-i+1}$

$$\begin{aligned} \sum_{i=1}^n \sum_{j>i} p_{ij} &= \sum_{i=1}^n \sum_{j>i} \frac{2}{j-i+1} \\ &\leq \sum_{i=1}^n \sum_{k=1}^{n-i+1} \frac{2}{k} \\ &\leq 2 \sum_{i=1}^n \sum_{k=1}^{n-i+1} \frac{1}{k} \\ &\leq 2nH_n \end{aligned}$$
- $H_n \approx \ln n + \Theta(1)$
- Therefore, **RandQS** $\in O(n \ln n)$


13.3 Analysis of Min-Cut (Monte Carlo)

13.3.1 RandMinCut Algorithm


- Let G be a connected, undirected multigraph (with multiple edges allowed between pairs of vertices).
- A *cut* in the graph is a set of edges whose removals from G result in G being broken into two or more components.
- The MINCUT is the set of such edges with smallest cardinality.
 - The following algorithm is a *randomized algorithm* for finding the minimum cut:
 1. Repeat the following until only two vertices remain:
 - * Remove an edge e uniformly at random
 - * Collapse (*contract*) the two vertex endpoints of e
 - All extra edges between the vertices collapsed are removed
 - All other edges are retained
 2. The edges that remain are a cut of G . But is it the minimum cut?

13.3.2 Example Walkthrough


1. Begin with the following graph




2. Select edge (b, e)




3. Now we selected edge $(b - e, c)$



4. Remove edge (f, d)



5. Finally, collapse $(b - e - c, a)$



6. The minimum cut of the graph G is $\{3, 6\}$

- This is clearly not the real minimum cut for G .
- **Practice:** What is the running time of this algorithm?

13.3.3 Analysis of RandMinCut

- Let k be the size of the actual minimum cut, and let \mathbf{C} be a cut of size k of a graph \mathbf{G} (whose number of vertices is at least $\frac{kn}{2}$).
- We will bound from below the probability that no edge of \mathbf{C} is ever removed during the execution of the algorithm, so that the set of edges left is a minimum cut.

1. Let ε_i denote the event that the i^{th} step of the algorithm does *not* remove an edge of \mathbf{C} , for $1 \leq i \leq n-2$. Note that n dictates the number of steps of the algorithm.

2. The probability that an edge of \mathbf{C} randomly chosen in the first step is *at most* $\frac{k}{\binom{nk}{2}} = \frac{2}{n}$.

– This implies $\Pr[\varepsilon_1] = 1 - \frac{2}{n}$

3. The probability that the edge randomly chosen in the second step is *at most* $\frac{2}{n-1}$.

– Therefore, $\Pr[\varepsilon_2 | \varepsilon_1] \geq 1 - \frac{2}{n-1}$

4. At the i^{th} step, there are $n-i+1$ vertices remaining, and the size of the minimum cut is still k , so the graph has at least $\frac{k(n-i+1)}{2}$ edges remaining.

5. Thus,

$$\Pr[\varepsilon_i | \cap_{j=1}^{i-1} \varepsilon_j] \geq 1 - \frac{2}{n-i+1}$$

6. What's the probability of never picking an edge of \mathbf{C} ?

- From our equation above, we have that

$$\Pr[\cap_{i=1}^n \varepsilon_i] = \Pr[\varepsilon_1] \times \Pr[\varepsilon_2 | \varepsilon_1] \times \Pr[\varepsilon_3 | \varepsilon_1 \cap \varepsilon_2] \dots \Pr[\varepsilon_k | \cap_{i=1}^{k-1} \varepsilon_i]$$

- This implies that

$$\Pr[\cap_{i=1}^{n-2} \varepsilon_i] \geq \prod_{i=1}^{n-1} \left(1 - \frac{2}{n-i+1}\right) = \frac{2}{n(n-1)}$$

7. Therefore, the probability of discovering a minimum cut is at least $\frac{2}{n^2}$, which means we may not (and probably won't) find the minimum cut for \mathbf{G} .

- Suppose we performed the same algorithm $\frac{n^2}{2}$ times, each independent.

– Recall:

$$\mathbf{Pr}[\varepsilon_1 \cap \varepsilon_2] = \mathbf{Pr}[\varepsilon_1] \times \mathbf{Pr}[\varepsilon_2]$$

– Therefore, the probability of *not* finding a minimum cut with $\frac{n^2}{2}$ runs of the algorithm, we get

$$\left(1 - \frac{2}{n^2}\right)^{\frac{n^2}{2}} < \frac{1}{\epsilon}$$

– Further executions of the algorithm allow for *arbitrarily small* failure probabilities ².

* Instead of $\frac{n^2}{2}$ executions, pick an arbitrarily small failure chance (ϵ) and solve for the number of executions necessary.

* For example, for 0.0001 failure chance, make episode 10000:

$$\begin{aligned} \left(1 - \frac{2}{n^2}\right)^X &< \frac{1}{10000} \\ X \lg\left(1 - \frac{2}{n^2}\right) &\leq \lg\left(\frac{1}{10000}\right) \\ X &\geq \frac{-\lg(10000)}{\lg\left(1 - \frac{2}{n^2}\right)} \end{aligned}$$

²This is the essence of Monte Carlo algorithms

13.4 Challenge Problem - Analysis of Vector Ordering

- We are posed the following question: Given an array A of length n of integers, is A in sorted order?
- We want a *sublinear time* algorithm for determining, with some failure probability, whether A is sorted.
- To accomplish this, we will use the following randomized algorithm:
 - Select a pair of elements at random
 - If they are in sorted order, return **true**. Otherwise, return **false**.
- Analyze the algorithm by answering the following questions:
 1. What is the total number of pairs from which to choose? Assume that two unique elements are chosen, so an element can't be paired with itself.
 2. Assume that m is the total number of *unsorted pairs* in A . That is, given the list of all of the pairs you may choose from, m is the total number of them that are not in sorted order.
 - What does it mean when $m = 0$?
 - What does it mean when $m = 1$?
 3. Let X_i be the i^{th} pair chosen, and let X be the total number of unsorted pairs found after n random pairs are drawn.
 - $X_i = 0$ when the i^{th} pair chosen is not sorted.
 - * In terms of n and m , what is the probability $X_i = 0$?
 - $X_i = 1$ when the i^{th} pair chosen is sorted.
 - * In terms of n and m , what is the probability $X_i = 1$?
 - What is the expected value of X ?
 4. What is an expression for the number of trials of the algorithm necessary to meet arbitrary success probability $\frac{1}{\epsilon}$, provided a value for m ? In otherwords, finish the phrase:
 - “If A has 5 unsorted pairs, I would need to select <some number> pairs to have a probability of $\frac{1}{\epsilon}$ of finding them.”